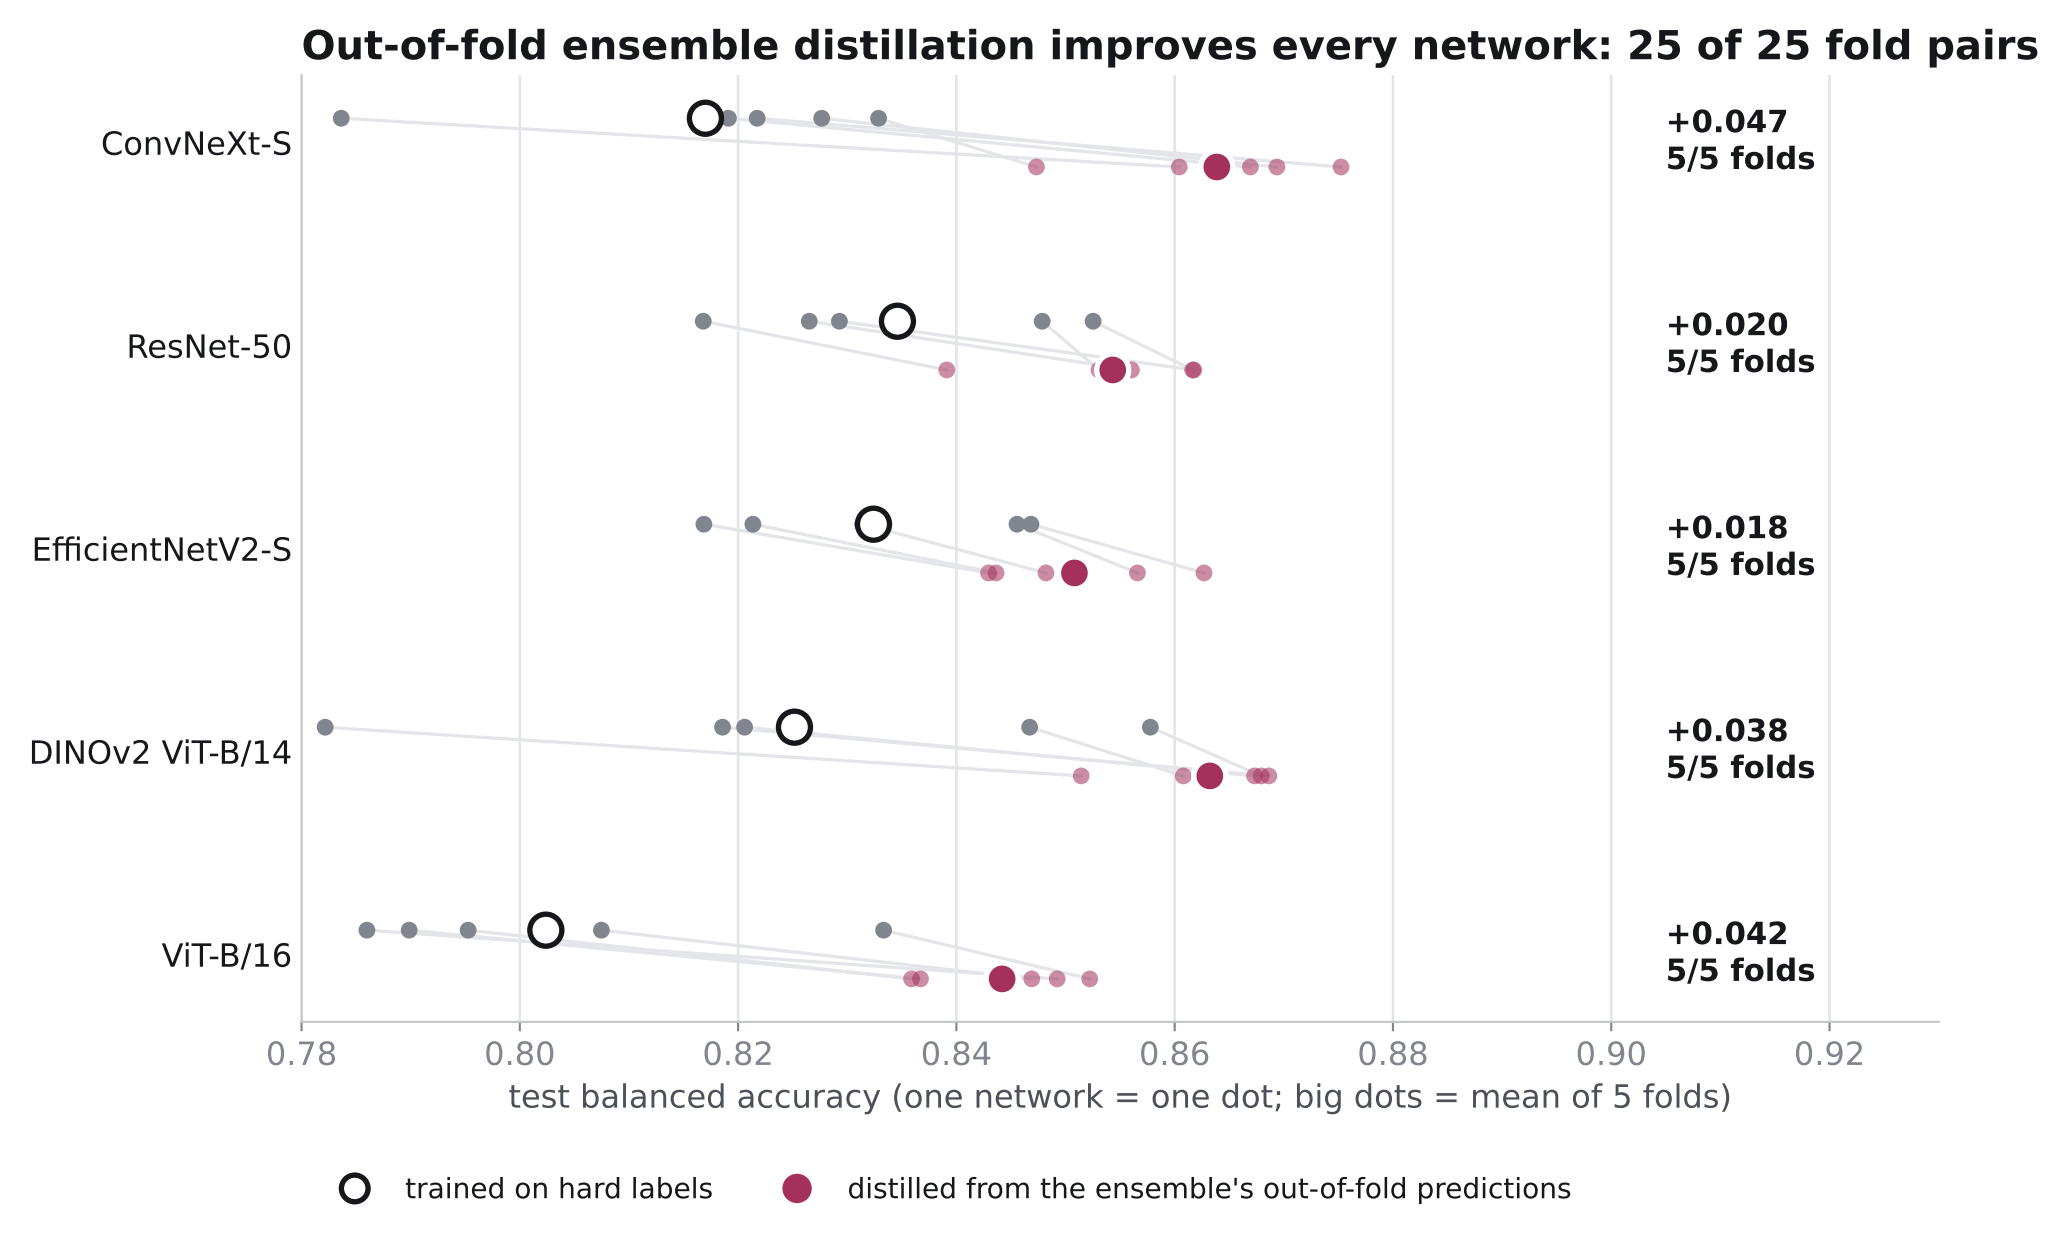
Out-of-fold ensemble distillation improves every network: 25 of 25 fold pairs
ConvNeXt-S
ResNet-50
EfficientNetV2-S
DINOv2 ViT-B/14
ViT-B/16
+0.047
5/5 folds
+0.020
5/5 folds
+0.018
5/5 folds
+0.038
5/5 folds
+0.042
5/5 folds
0.78
0.80
0.82
0.84
0.86
0.88
0.90
0.92
test balanced accuracy (one network = one dot; big dots = mean of 5 folds)
trained on hard labels
distilled from the ensemble's out-of-fold predictions