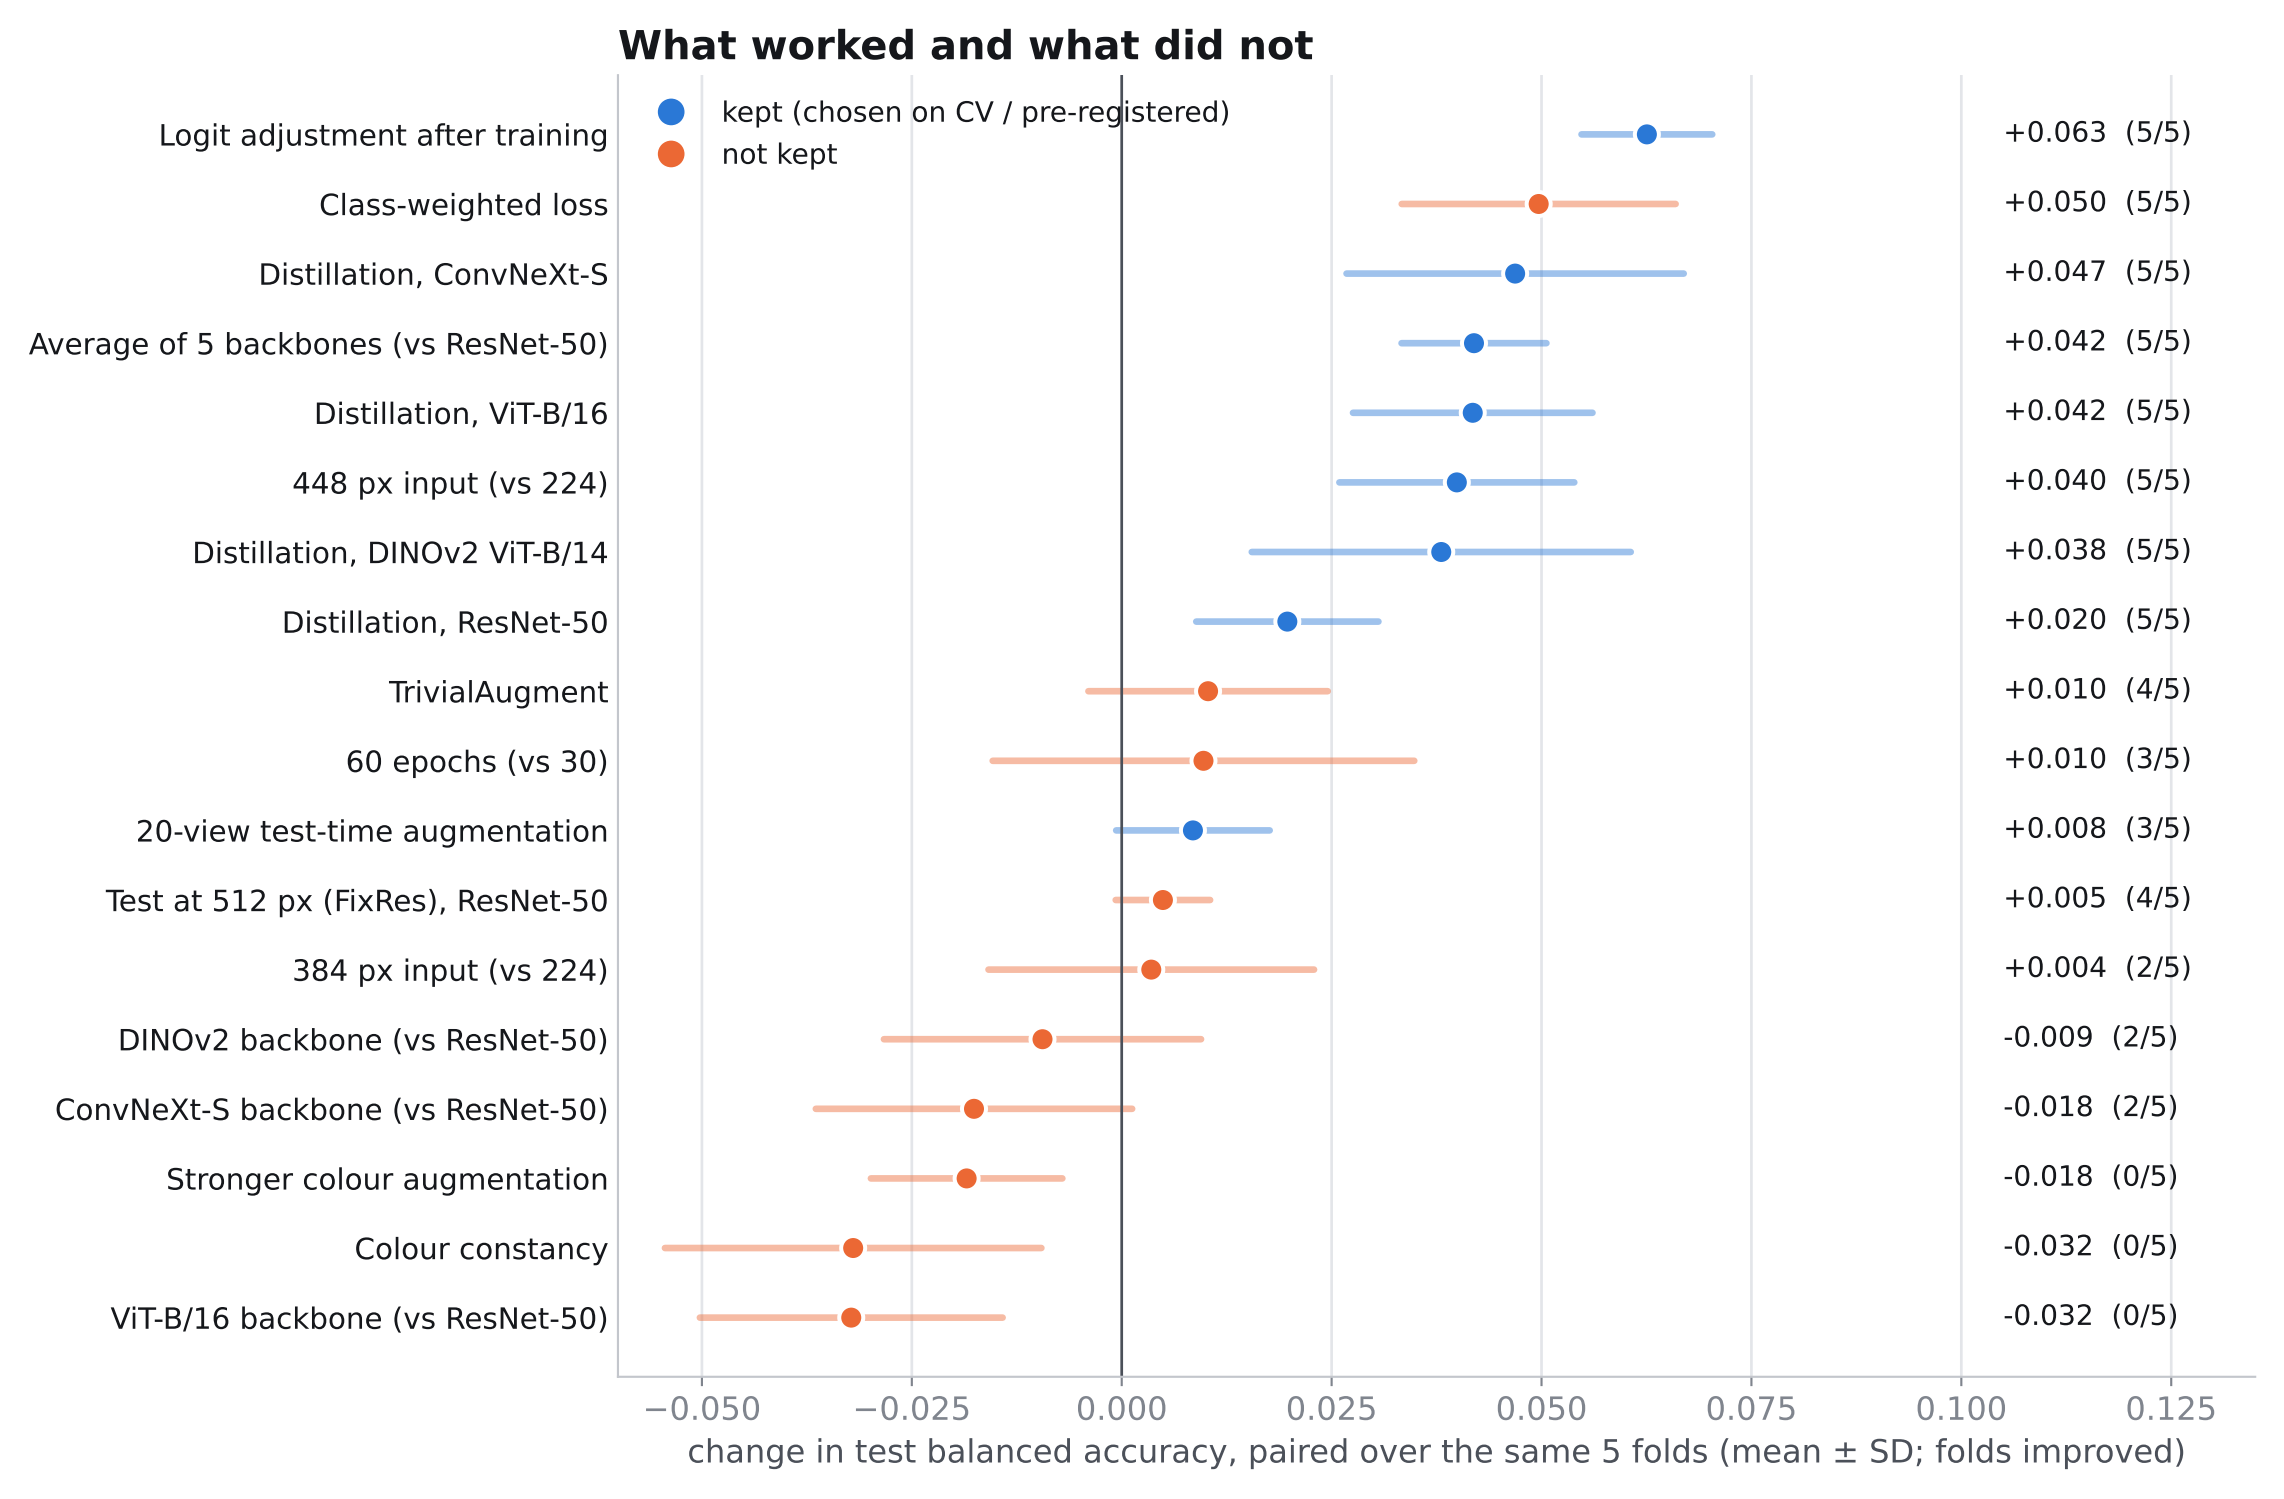
What worked and what did not
kept (chosen on CV / pre-registered)
not kept
Logit adjustment after training
Class-weighted loss
Distillation, ConvNeXt-S
Average of 5 backbones (vs ResNet-50)
Distillation, ViT-B/16
448 px input (vs 224)
Distillation, DINOv2 ViT-B/14
Distillation, ResNet-50
TrivialAugment
60 epochs (vs 30)
20-view test-time augmentation
Test at 512 px (FixRes), ResNet-50
384 px input (vs 224)
DINOv2 backbone (vs ResNet-50)
ConvNeXt-S backbone (vs ResNet-50)
Stronger colour augmentation
Colour constancy
ViT-B/16 backbone (vs ResNet-50)
+0.063 (5/5)
+0.050 (5/5)
+0.047 (5/5)
+0.042 (5/5)
+0.042 (5/5)
+0.040 (5/5)
+0.038 (5/5)
+0.020 (5/5)
+0.010 (4/5)
+0.010 (3/5)
+0.008 (3/5)
+0.005 (4/5)
+0.004 (2/5)
-0.009 (2/5)
-0.018 (2/5)
-0.018 (0/5)
-0.032 (0/5)
-0.032 (0/5)
−0.050
−0.025
0.000
0.025
0.050
0.075
0.100
0.125
change in test balanced accuracy, paired over the same 5 folds (mean ± SD; folds improved)