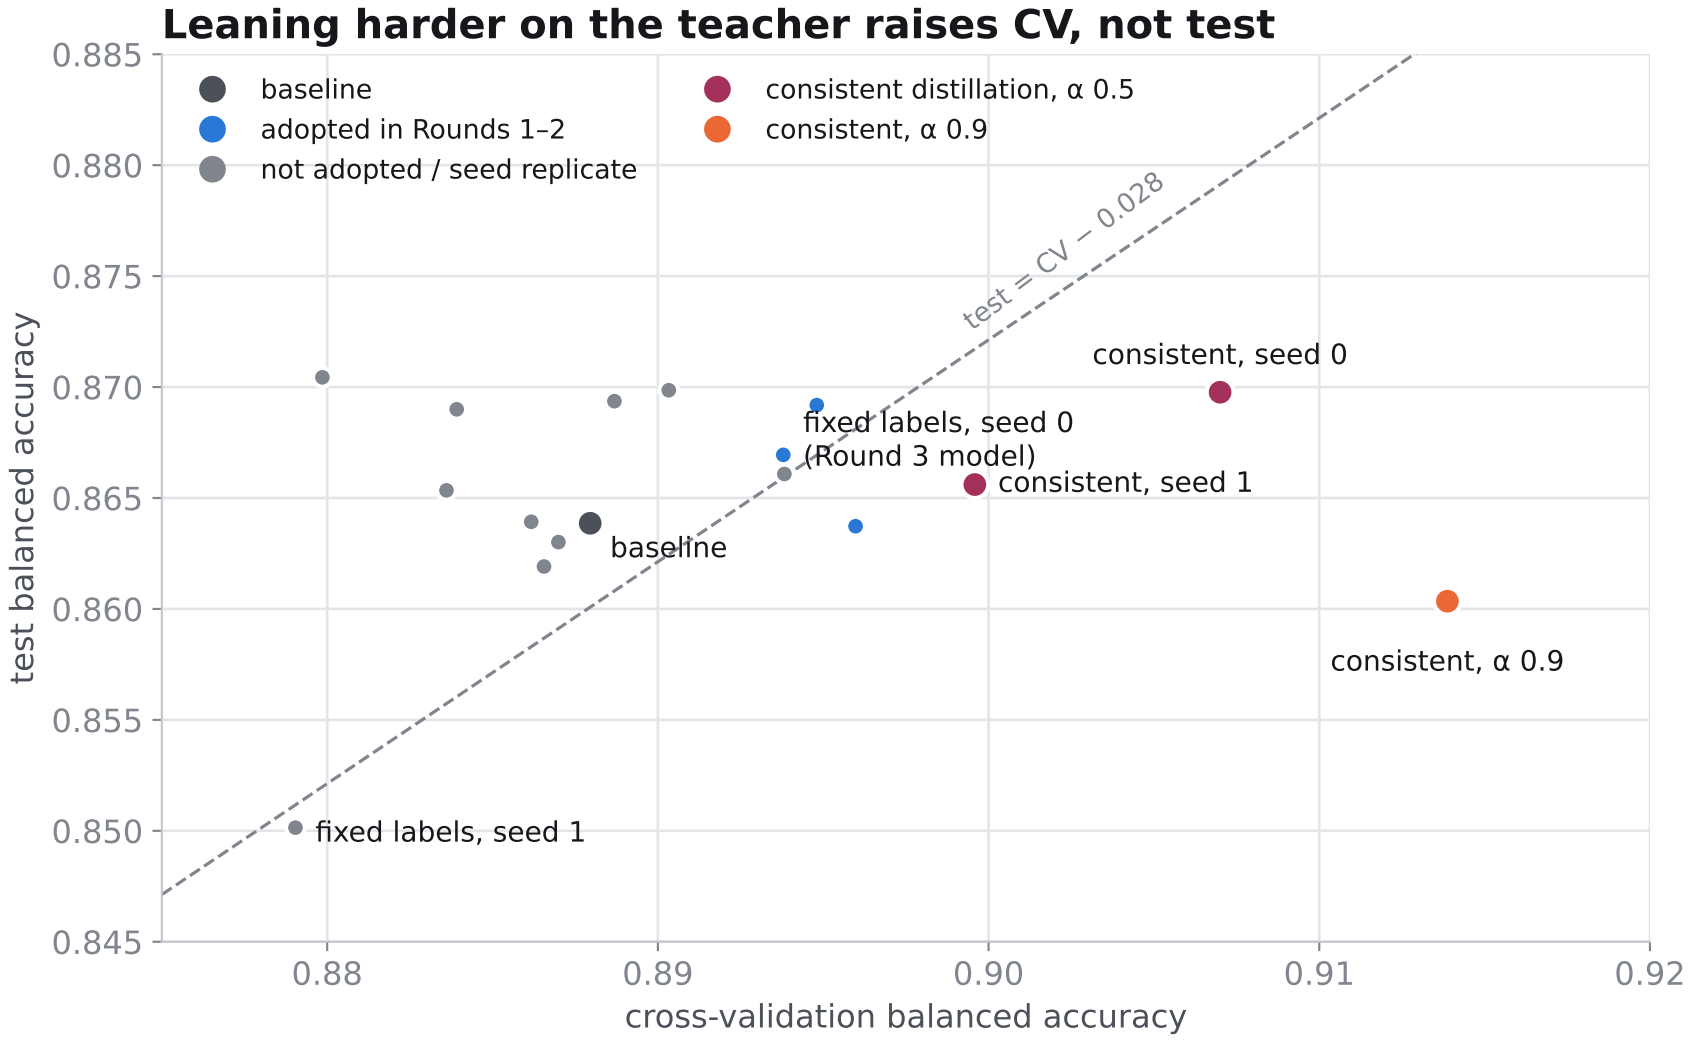

Each dot is one network (5-fold mean); dashed line: the CV − test gap of the fixed-label runs. The teacher was trained on the student's CV images, so copying it more (α 0.9) flatters CV: its CV − test gap is 0.054, twice the 0.028 of the fixed-label runs.